



Structured Query Language

Chapter 5

Contents

-
- 1 Introduction of Structured Query Language
 - 2 DDL: create, drop, alter
 - 3 DML: select, insert, update, delete
 - 4 DCL: commit, rollback, grant, revoke

Contents

- 1 Introduction of Structured Query Language**
- 2 DDL: create, drop, alter
- 3 DML: select, insert, update, delete
- 4 DCL: commit, rollback, grant, revoke

Introduction of Structured Query Language

- ▶ **Structured Query Language (SQL)** is a standard computer language for relational database management and data manipulation.
- ▶ **Basic SQL:**
 - ▶ **Data Definition Language (DDL)**
 - ▶ Create, Alter, Drop
 - ▶ **Data Manipulation Language (DML)**
 - ▶ Select, Insert, Update, Delete
 - ▶ **Data Control Language (DCL)**
 - ▶ Commit, Rollback, Grant, Revoke

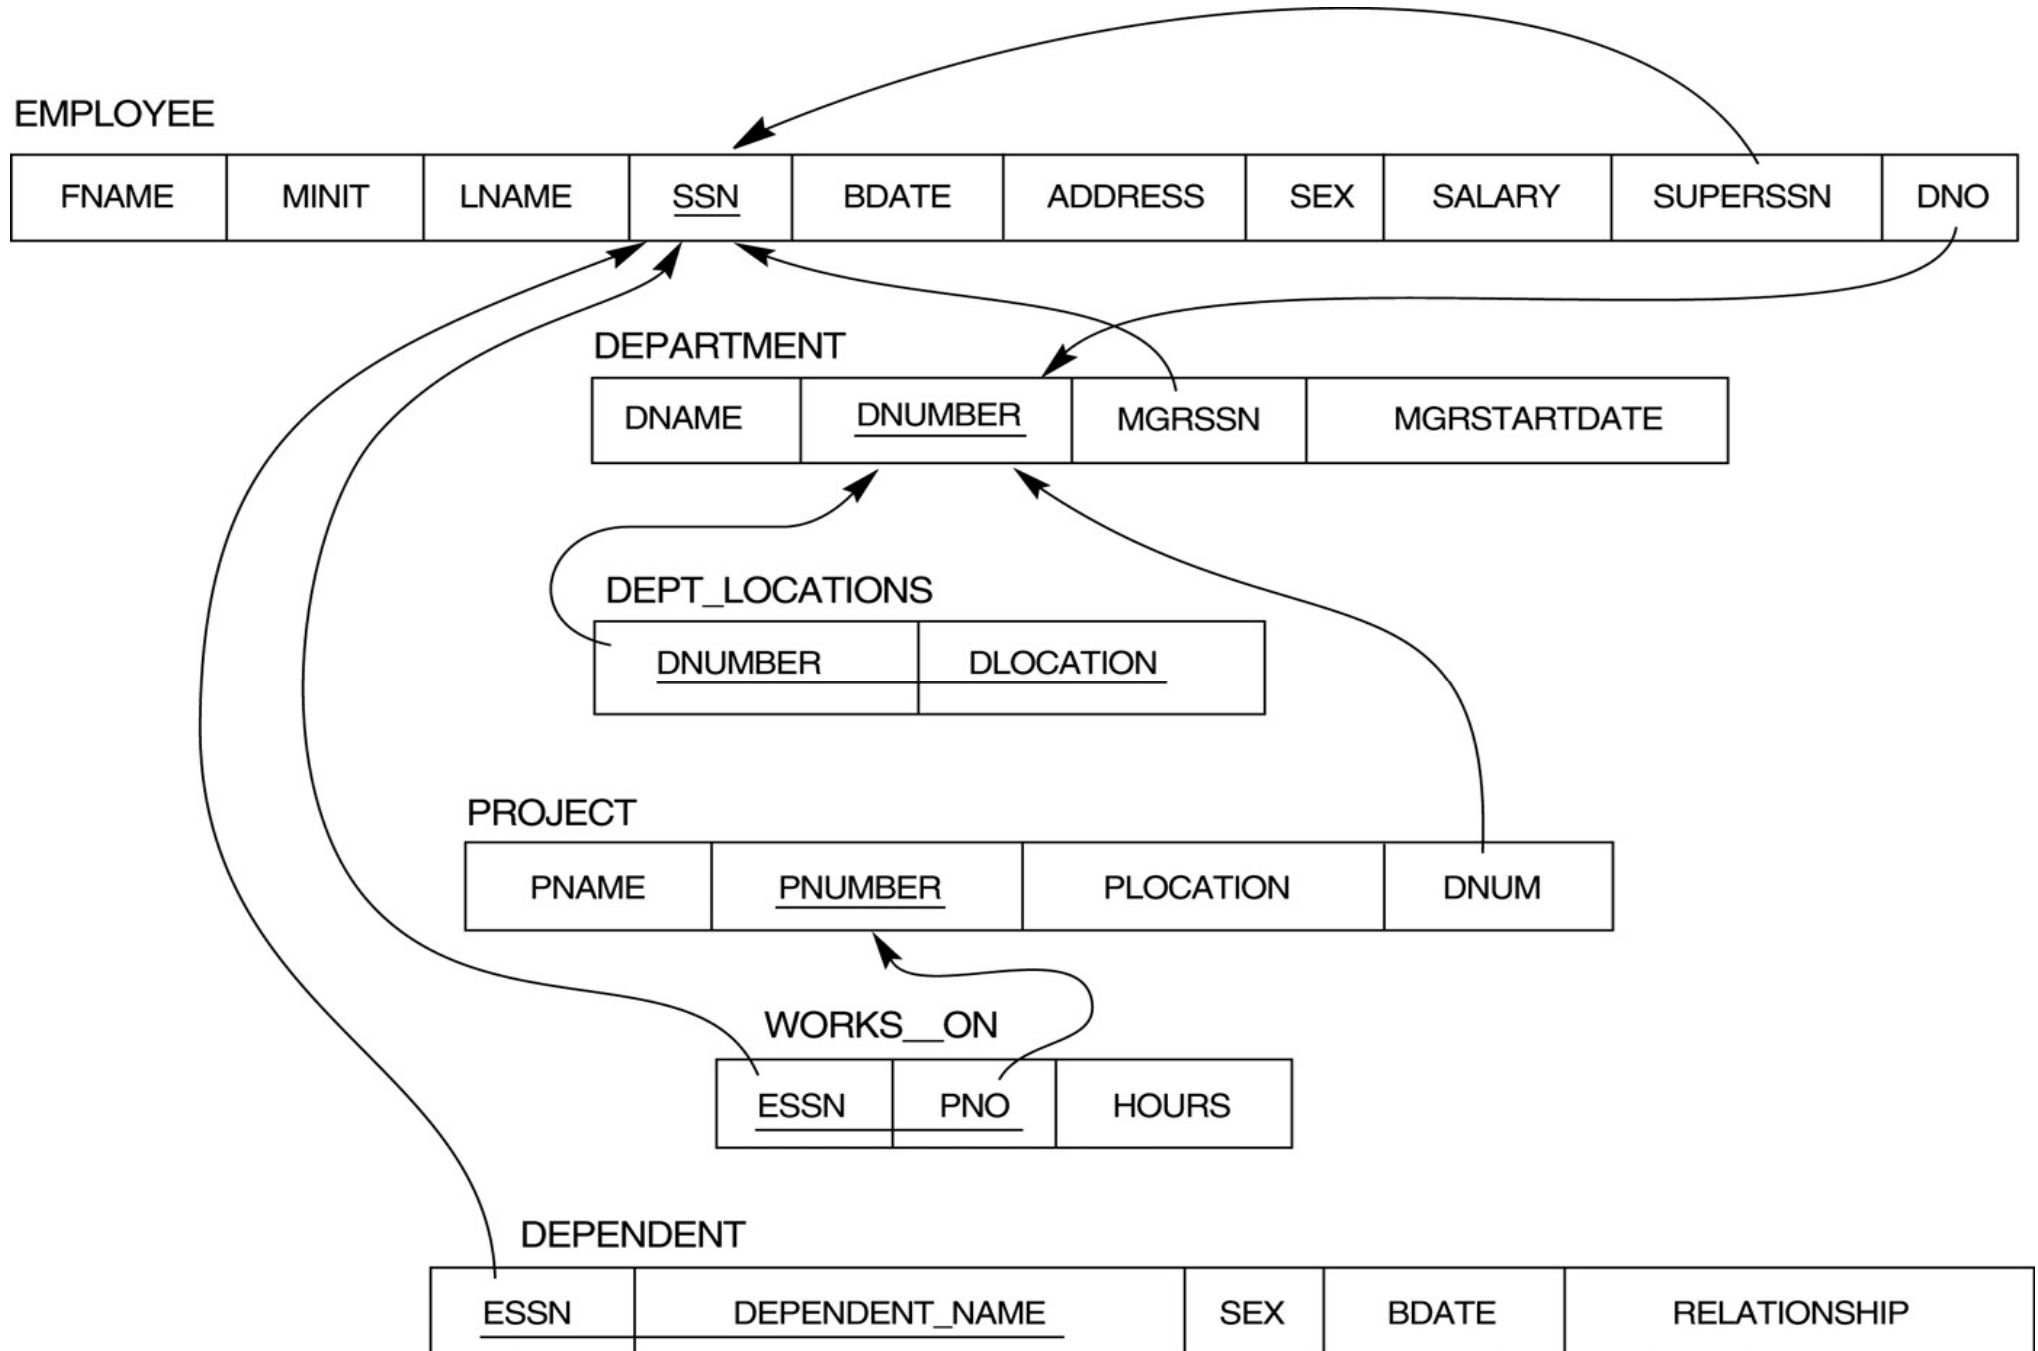
Contents

- 1 Introduction of Structured Query Language
- 2 DDL: create, drop, alter**
- 3 DML: select, insert, update, delete
- 4 DCL: commit, rollback, grant, revoke

Data Definition Language (DDL)

- ▶ Permits specification of data types, structures and any data constraints
- ▶ All specifications are stored in the database
- ▶ Includes:
 - ▶ **CREATE**: make a new database object (database, table, index, user, stored query, ...)
 - ▶ **ALTER**: modify an existing database object
 - ▶ **DROP**: destroy an existing database object

The COMPANY Database



Schema and Catalog Concepts in SQL

- ▶ **Schema:** a group of tables and other constructs that belong to the same database application

```
CREATE SCHEMA Schema_Name  
AUTHORIZATION Authorization_Identifier;
```

```
CREATE SCHEMA Company AUTHORIZATION JSmith;
```

- ▶ **Catalog:** a named collection of schemas

CREATE TABLE

```
CREATE TABLE [SchemaName.]TableName
({colName dataType [NOT NULL] [UNIQUE] [PRIMARY KEY]
[DEFAULT defaultOption]
[CHECK searchCondition] [...]}
[PRIMARY KEY (listOfColumns),]
{[UNIQUE (listOfColumns),] [...]}
{[FOREIGN KEY (listOfFKColumns)
REFERENCES ParentTableName [(listOfCKColumns)]
[ON UPDATE referentialAction]
[ON DELETE referentialAction ]] [...]}
{[CHECK (searchCondition)] [...] })
```

CREATE TABLE

- ▶ **Base tables** (base relations)
 - ▶ Relation and its tuples are actually created and stored as a file by the DBMS.
- ▶ **Virtual relations**
 - ▶ Created through the **CREATE VIEW** statement.
- ▶ Some foreign keys may cause errors
 - ▶ Circular references
 - ▶ refer to a table that has not yet been created

Basic Data Types

- ▶ **Numeric data types:**

- ▶ Integer numbers: INTEGER, INT, and SMALLINT
- ▶ Floating-point (real) numbers: FLOAT or REAL, and DOUBLE PRECISION

- ▶ **Character-string data types:**

- ▶ Fixed length: CHAR(n), CHARACTER(n)
- ▶ Varying length: VARCHAR(n), CHAR VARYING(n), CHARACTER VARYING(n)

- ▶ **Bit-string data types:**

- ▶ Fixed length: BIT(n)
- ▶ Varying length: BIT VARYING(n)

- ▶ **Boolean data type:**

- ▶ Values of TRUE or FALSE or NULL

- ▶ **Date-Time data types:**

- ▶ Date components: YEAR, MONTH, and DAY ('YYYY-MM-DD')
- ▶ Time components: HOUR, MINUTE, and SECOND ('HH:MM:SS')

Basic Data Types

- ▶ Additional data types
 - ▶ **Timestamp data type (TIMESTAMP)**
 - ▶ Includes the DATE and TIME fields
 - ▶ Plus a minimum of six positions for decimal fractions of seconds
 - ▶ Optional WITH TIME ZONE qualifier
 - ▶ **INTERVAL data type**
 - ▶ Specifies a relative value that can be used to increment or decrement an absolute value of a date, time, or timestamp

Domains

- ▶ Name used with the attribute specification
- ▶ Makes it easier to change the data type for a domain that is used by numerous attributes
- ▶ Improves schema readability

```
CREATE DOMAIN DomainName AS Data Type  
[CHECK conditions];
```

```
CREATE DOMAIN SSN_TYPE AS CHAR(9);  
CREATE DOMAIN D_NUM AS INTEGER  
CHECK (D_NUM>0 AND D_NUM<21);
```

Specifying Constraints

- ▶ Basic constraints:
 - ▶ Key and referential integrity constraints
 - ▶ Attribute constraints
 - ▶ Constraints on individual tuples within a relation

Key and Referential Integrity Constraints

- ▶ **PRIMARY KEY** clause: specifies one or more attributes that make up the primary key of a relation.

Dnumber INT **PRIMARY KEY**

PRIMARY KEY (Dnumber, DLocation)

- ▶ **UNIQUE** clause: Specifies alternate (secondary) keys.

Dname VARCHAR(15) **UNIQUE**;

Key and Referential Integrity Constraints

► FOREIGN KEY clause

FOREIGN KEY (*listOfFKColumns*)
REFERENCES *ParentTableName* [(*listOfCKColumns*)]
[**ON UPDATE** *referentialAction*]
[**ON DELETE** *referentialAction*]

- Referential triggered actions: RESTRICT (default), SET NULL, CASCADE, and SET DEFAULT

FOREIGN KEY Dno **REFERENCES** Department(Dnumber)
ON DELETE CASCADE
ON UPDATE CASCADE

Attribute Constraints

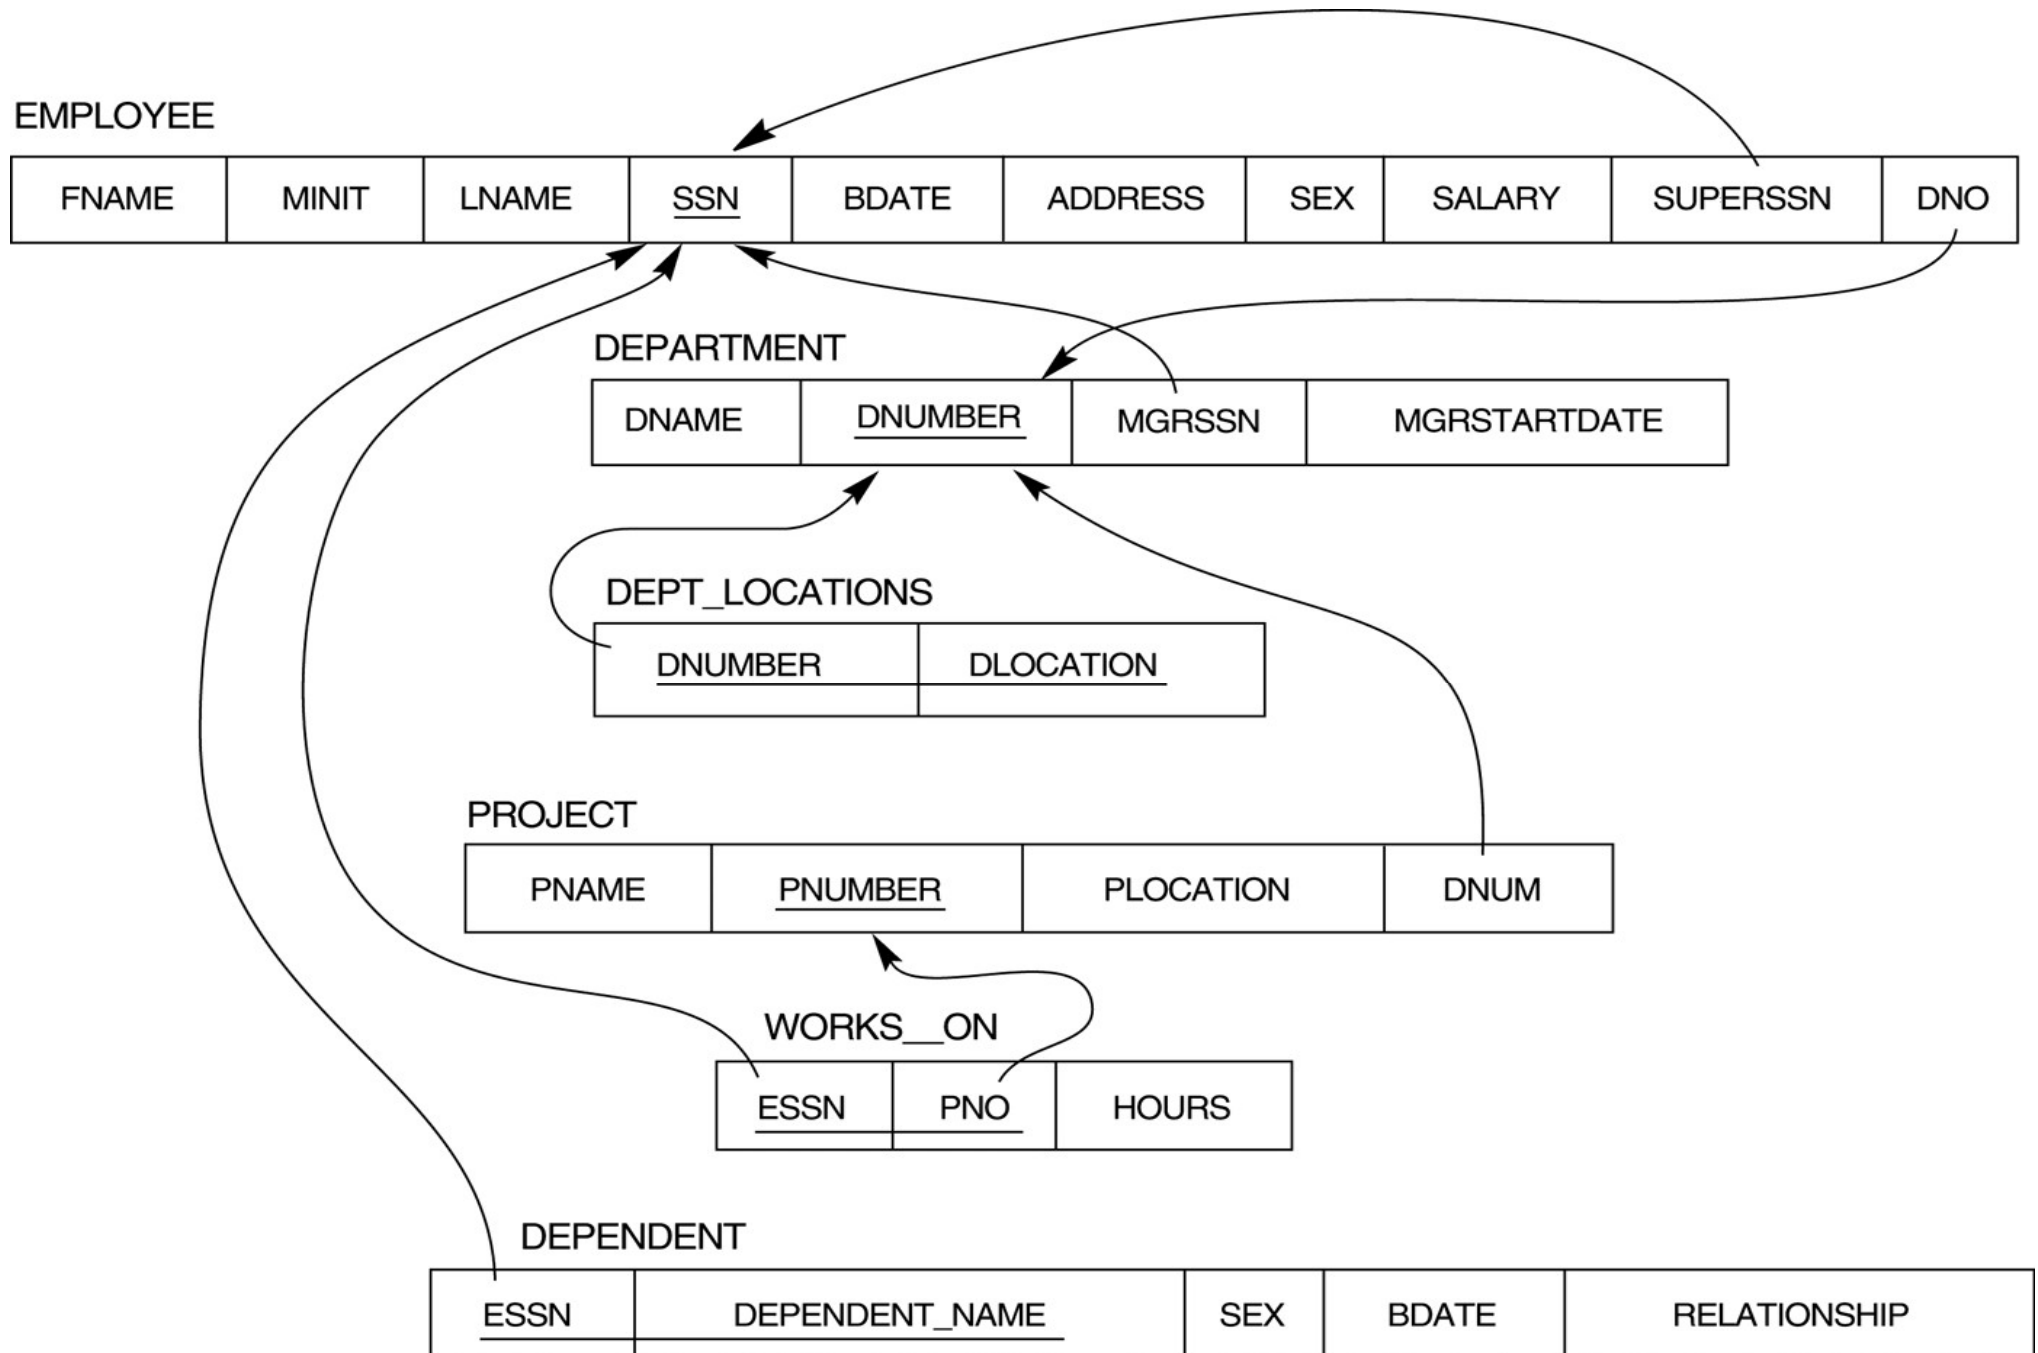
- ▶ NOT NULL
 - ▶ NULL is **not** permitted for a particular attribute
- ▶ Default values
 - ▶ DEFAULT <value> can be specified for an attribute
 - ▶ If no default clause is specified, the default value is NULL for attributes that do not have the NOT NULL constraint

Dno INT NOT NULL DEFAULT 1

- ▶ CHECK clause:

Dnumber INT NOT NULL CHECK (Dnumber > 0 AND Dnumber <21);

The COMPANY Database



CREATE TABLE EMPLOYEE

(Fname	VARCHAR(15)	NOT NULL,
Minit	CHAR,	
Lname	VARCHAR(15)	NOT NULL,
Ssn	CHAR(9)	NOT NULL,
Bdate	DATE,	
Address	VARCHAR(30),	
Sex	CHAR,	
Salary	DECIMAL(10,2),	
Super_ssn	CHAR(9),	
Dno	INT	NOT NULL,

PRIMARY KEY (Ssn),

CREATE TABLE DEPARTMENT

(Dname	VARCHAR(15)	NOT NULL,
Dnumber	INT	NOT NULL,
Mgr_ssn	CHAR(9)	NOT NULL,
Mgr_start_date	DATE,	

PRIMARY KEY (Dnumber),

UNIQUE (Dname),

FOREIGN KEY (Mgr_ssn) REFERENCES EMPLOYEE(Ssn));

CREATE TABLE DEPT_LOCATIONS

(Dnumber	INT	NOT NULL,
Dlocation	VARCHAR(15)	NOT NULL,

PRIMARY KEY (Dnumber, Dlocation),

FOREIGN KEY (Dnumber) REFERENCES DEPARTMENT(Dnumber));

CREATE TABLE PROJECT

(Pname	VARCHAR(15)	NOT NULL,
Pnumber	INT	NOT NULL,
Plocation	VARCHAR(15),	
Dnum	INT	NOT NULL,
PRIMARY KEY (Pnumber),		
UNIQUE (Pname),		
FOREIGN KEY (Dnum) REFERENCES DEPARTMENT(Dnumber));		

CREATE TABLE WORKS_ON

(Essn	CHAR(9)	NOT NULL,
Pno	INT	NOT NULL,
Hours	DECIMAL(3,1)	NOT NULL,
PRIMARY KEY (Essn, Pno),		
FOREIGN KEY (Essn) REFERENCES EMPLOYEE(Ssn),		
FOREIGN KEY (Pno) REFERENCES PROJECT(Pnumber));		

CREATE TABLE DEPENDENT

(Essn	CHAR(9)	NOT NULL,
Dependent_name	VARCHAR(15)	NOT NULL,
Sex	CHAR,	
Bdate	DATE,	
Relationship	VARCHAR(8),	
PRIMARY KEY (Essn, Dependent_name),		
FOREIGN KEY (Essn) REFERENCES EMPLOYEE(Ssn));		

Specifying Constraints

- ▶ Giving names to constraints
 - ▶ This is optional.
 - ▶ Keyword **CONSTRAINT**
 - ▶ The name is unique within a particular DB schema.
 - ▶ Used to identify a particular constraint in case it must be dropped later and replaced with another one.

CREATE TABLE EMPLOYEE

```
( ... ,  
  Dno          INT          NOT NULL          DEFAULT 1,  
  CONSTRAINT EMPPK  
    PRIMARY KEY (Ssn),  
  CONSTRAINT EMPSUPERFK  
    FOREIGN KEY (Super_ssn) REFERENCES EMPLOYEE(Ssn)  
      ON DELETE SET NULL          ON UPDATE CASCADE,  
  CONSTRAINT EMPDEPTFK  
    FOREIGN KEY(Dno) REFERENCES DEPARTMENT(Dnumber)  
      ON DELETE SET DEFAULT      ON UPDATE CASCADE);
```

CREATE TABLE DEPARTMENT

```
( ... ,  
  Mgr_ssn CHAR(9)          NOT NULL          DEFAULT '888665555',  
  ... ,  
  CONSTRAINT DEPTPK  
    PRIMARY KEY(Dnumber),  
  CONSTRAINT DEPTSK  
    UNIQUE (Dname),  
  CONSTRAINT DEPTMGRFK  
    FOREIGN KEY (Mgr_ssn) REFERENCES EMPLOYEE(Ssn)  
      ON DELETE SET DEFAULT      ON UPDATE CASCADE);
```

CREATE TABLE DEPT_LOCATIONS

```
( ... ,  
  PRIMARY KEY (Dnumber, Dlocation),  
  FOREIGN KEY (Dnumber) REFERENCES DEPARTMENT(Dnumber)  
    ON DELETE CASCADE          ON UPDATE CASCADE);
```

Constraints on individual tuples within a relation

- ▶ Specifying constraints on tuples using CHECK
 - ▶ Affected on each tuple individually as being inserted or modified (tuple-based constraints)
 - ▶ Ex: Department's create-date must be earlier than the manager's start-date:

CHECK (DEPT_CREATE_DATE < MGRSTARTDATE);

- ▶ More general constraints: CREATE ASSERTION

DROP Command

- ▶ Used to drop named schema elements: tables, domains, constraints, and the schema itself
- ▶ Drop behavior options:
 - ▶ CASCADE and RESTRICT

DROP SCHEMA Company CASCADE;

Or

DROP SCHEMA Company RESTRICT;

DROP Command

- ▶ Drop a table:

DROP TABLE Department **CASCADE**;

- ▶ RESTRICT (default): dropped on if it is not referenced in any constraints or views
- ▶ CASCADE: all such constraints and views that reference the table are dropped automatically from the schema along with the table itself
- ▶ Similarly, we can drop constraints & domains

ALTER Command

- ▶ ALTER command: change the definition of a base table or of other named schema elements
- ▶ Base tables: adding or dropping a column or constraints, changing a column definition.

ALTER TABLE Employee ADD Job VARCHAR(15);

ALTER TABLE Employee

DROP COLUMN Address CASCADE;

ALTER TABLE Department

ALTER COLUMN Mgr_ssn SET DEFAULT '333445555';

ALTER TABLE Employee

DROP CONSTRAINT Empsuperfk CASCADE;

ALTER TABLE Employee

Foreign key Dno references Department(Dnumber);

Contents

-
- 1 Introduction of Structured Query Language
 - 2 DDL: create, drop, alter
 - 3 DML: select, insert, update, delete**
 - 4 DCL: commit, rollback, grant, revoke

SELECT Command

- ▶ SELECT command: retrieve information from a database
- ▶ SELECT command in SQL is the same as the SELECT operation in relational algebra.
- ▶ SQL allows a table (relation) to have two or more tuples that are identical in all their attribute values
- ▶ SQL relation (table) is a multi-set (sometimes called a bag) of tuples; it is not a set of tuples
- ▶ SQL relations can be constrained to be sets by specifying PRIMARY KEY or UNIQUE attributes, or by using the DISTINCT option in a query

SELECT Command

- ▶ Basic form:

SELECT *<attribute list>*
FROM *<table list>*
WHERE *<condition>*

- ▶ *<attribute list>* is a list of attribute names whose values are to be retrieved by the query
- ▶ *<table list>* is a list of the relation names required to process the query
- ▶ *<condition>* is a conditional (Boolean) expression that identifies the tuples to be retrieved by the query

SELECT Command

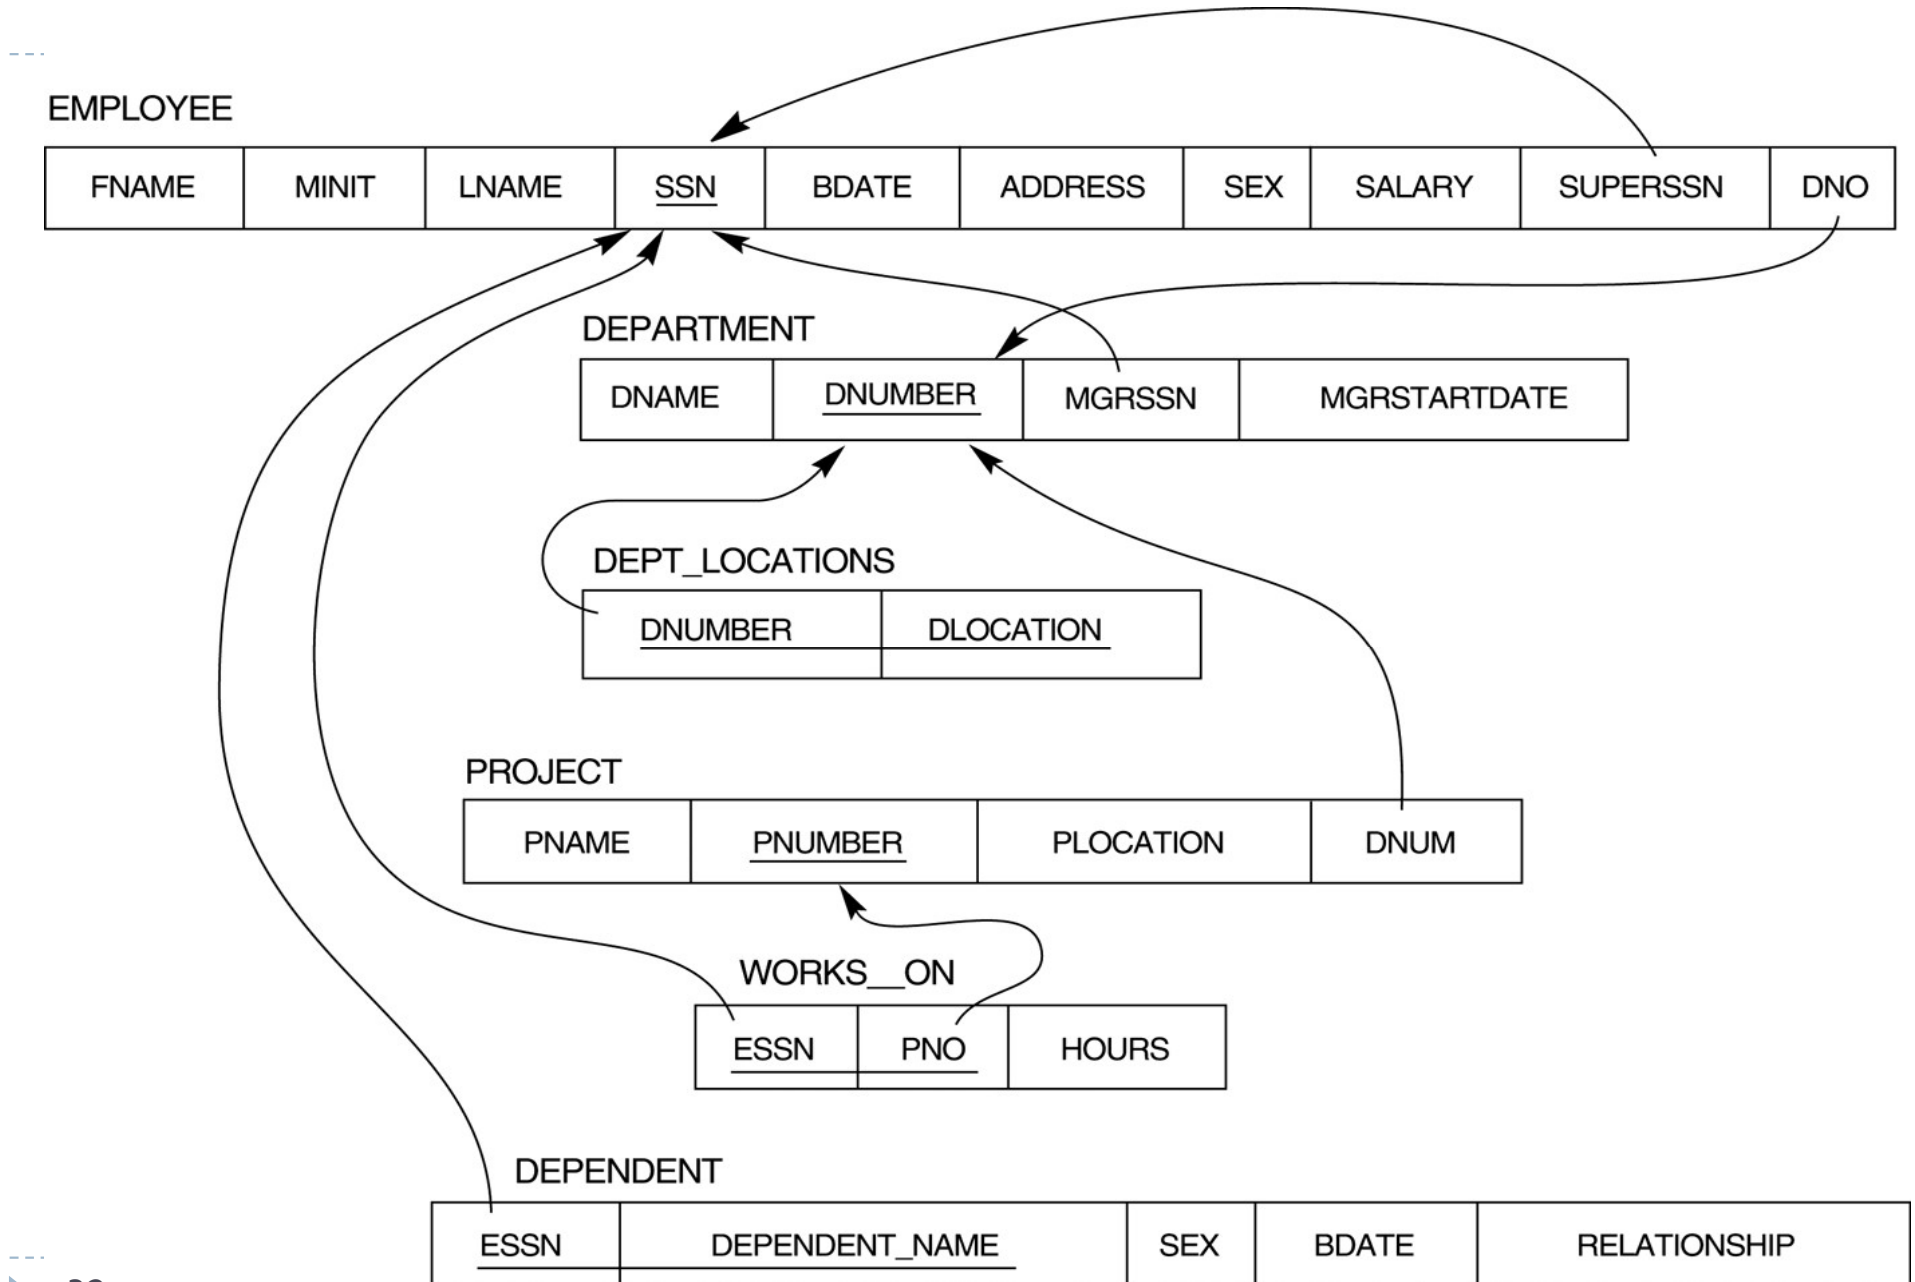
- ▶ Complete form:

```
SELECT [DISTINCT | ALL]
      { * | [columnExpression [AS newName]] [,...] }
FROM TableName [alias] [, ...]
[WHERE condition]
[GROUP BY columnList]    [HAVING condition]
[ORDER BY columnList]
```

SELECT Command

- ▶ **SELECT** : Specifies which columns are to appear in output
- ▶ **FROM** : Specifies table(s) to be used
- ▶ **WHERE** : Filters rows
- ▶ **GROUP BY** : Forms groups of rows with same column value
- ▶ **HAVING** : Filters groups subject to some condition
- ▶ **ORDER BY** : Specifies the order of the output

The COMPANY Database



SELECT Command

- ▶ Basic SQL queries: using the SELECT, PROJECT, and JOIN operations of the relational algebra

Query 0: Retrieve the birthdate and address of the employee whose name is 'John B. Smith'.

Q0: **SELECT** Bdate, Address
 FROM Employee
 WHERE Fname = 'John' **AND** Minit = 'B'
 AND Lname = 'Smith';

- ▶ Similar to a SELECT-PROJECT pair of relational algebra operations:
 - ▶ SELECT clause specifies the projection attributes
 - ▶ WHERE clause specifies the selection condition
 - ▶ However, the result of the query may contain **duplicate tuples**

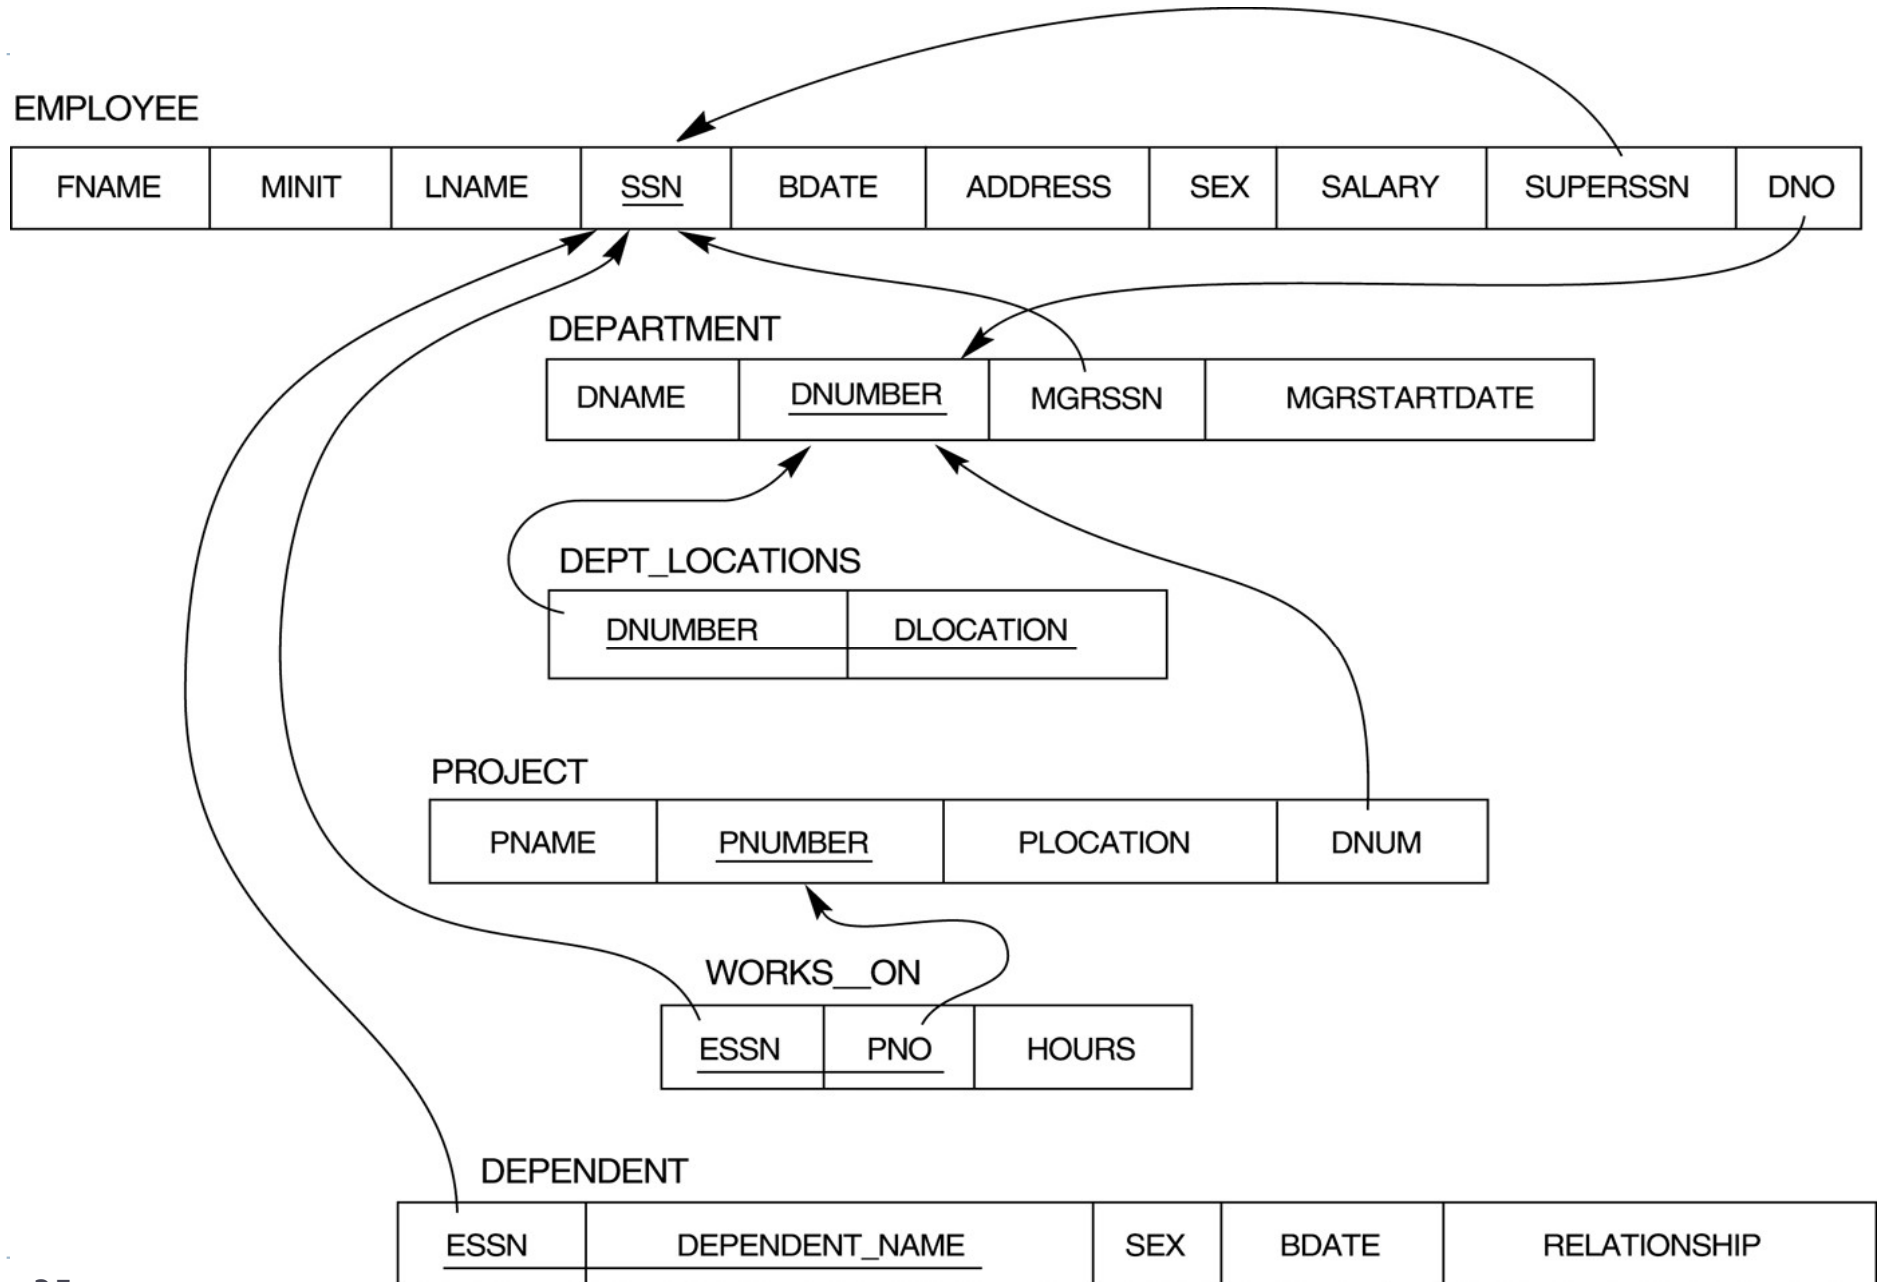
SELECT Command

Query 1: Retrieve the name and address of all employees who work for the 'Research' department.

Q1: **SELECT** Fname, Lname, Address
 FROM Employee, Department
 WHERE Dname='Research' **AND** Dnumber= Dno;

- ▶ Similar to a SELECT-PROJECT-JOIN sequence of relational algebra operations
 - ▶ (DNAME='Research'): selection condition (SELECT operation in relational algebra)
 - ▶ (DNUMBER=DNO): join condition (JOIN operation in relational algebra)

The COMPANY Database



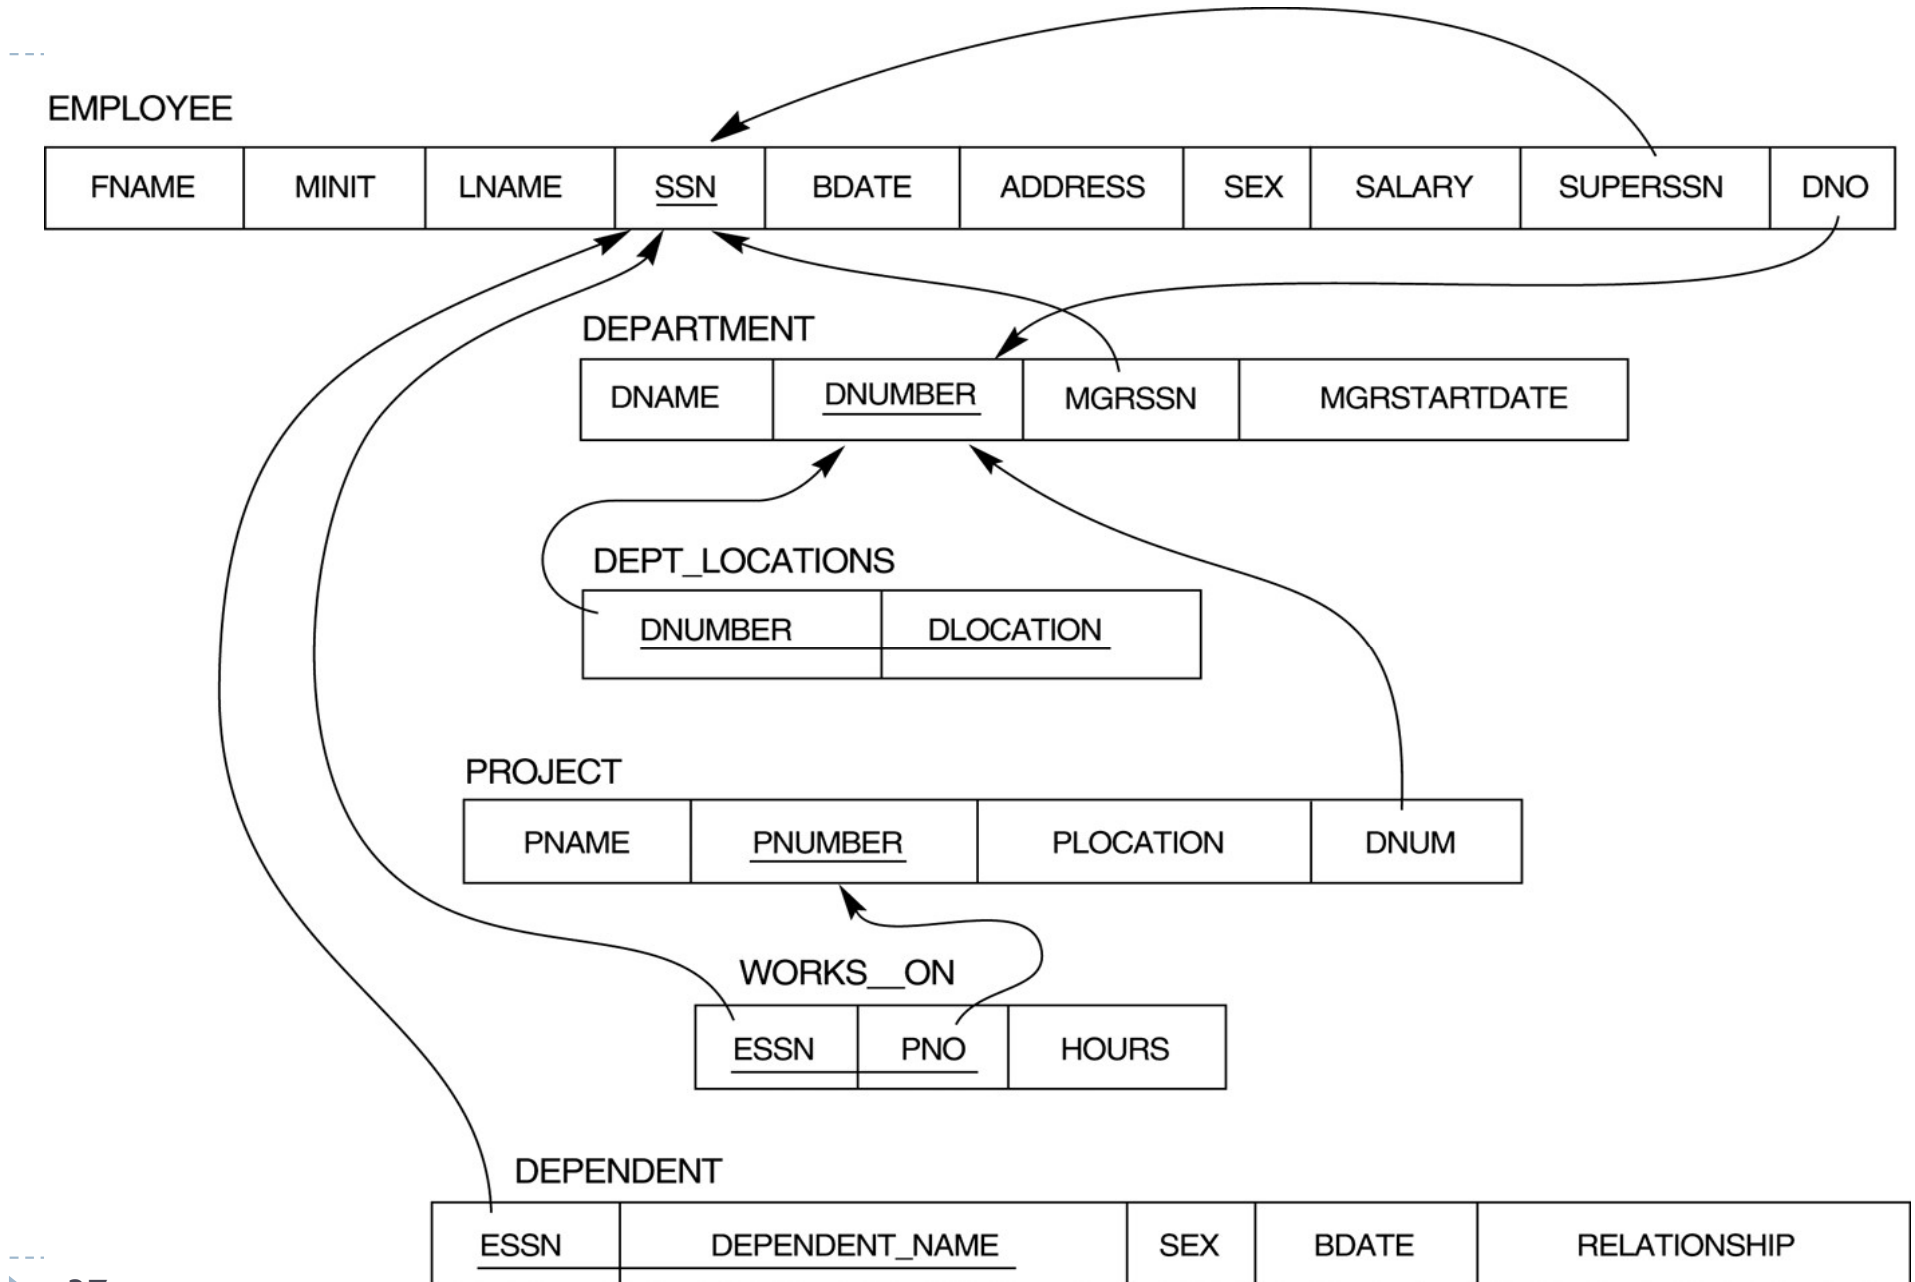
SELECT Command

Query 2: For every project located in 'Stafford', list the project number, the controlling department number, and the department manager's last name, address, and birthdate

Q2: **SELECT** Pnumber, Dnum, Lname, Bdate, Address
 FROM Project, Department, Employee
 WHERE Dnum = Dnumber **AND** MgrSSN = SSN
 AND Plocation='Stafford';

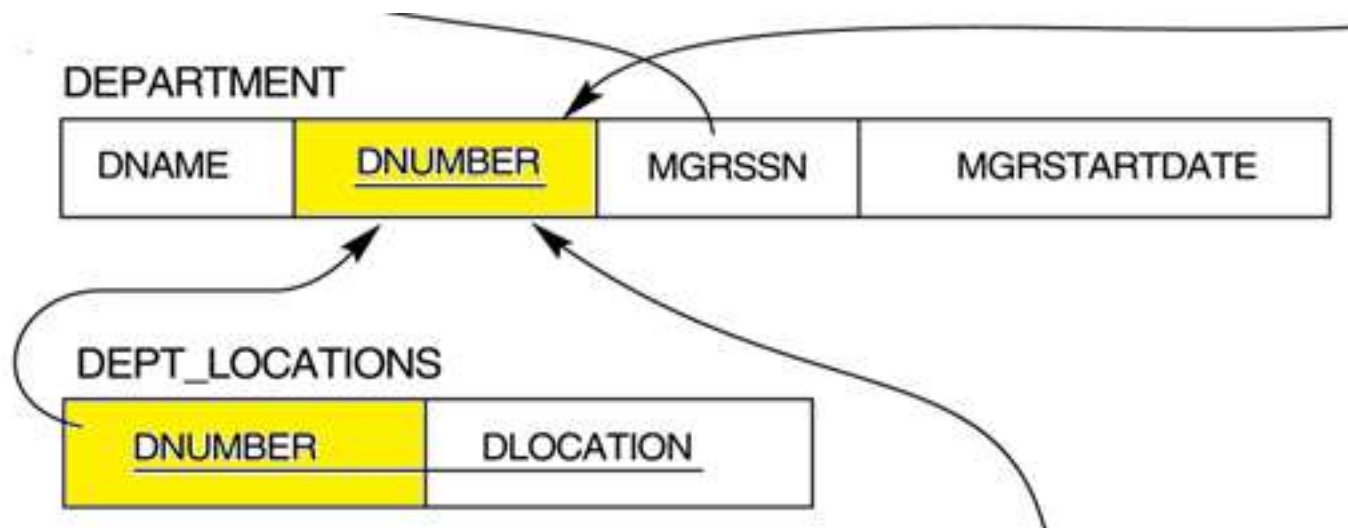
- ▶ Two join conditions:
 - ▶ Dnum = Dnumber: relates a project to its controlling department
 - ▶ MgrSSN = SSN: relates the controlling department to the employee who manages that department

The COMPANY Database



Ambiguous Attribute Names

- ▶ In SQL, we can use the same name for attributes as long as the attributes are in *different relations*. Query referring to attributes with the same name must *qualify* the attribute name with the relation name by *prefixing* the relation name to the attribute name
- ▶ Examples:
 - ▶ DEPARTMENT.DNUMBER and DEPT_LOCATIONS.DNUMBER



Aliases

- ▶ Some queries need to refer to the same relation twice: aliases are given to the relation name

Query 3: For each employee, retrieve the employee's name, and the name of his or her immediate supervisor.

Q3a: **SELECT** E.Fname, E.Lname, S.Fname, S.Lname
FROM Employee E S
WHERE E.SuperSSN = S.SSN;

- ▶ E and S are called aliases or tuple variables for the Employee relation
 - ▶ E represents employees in role of supervisees
 - ▶ S represents employees in role of supervisors

Aliases

- ▶ Aliases can also be used in any SQL query for convenience. Can also use the AS keyword to specify aliases

Q3b: **SELECT** E.Fname, E.Lname, S.Fname, S.Lname
FROM Employee **AS** E, Employee **AS** S
WHERE E.SuperSSN = S.SSN;

- ▶ Renaming using aliases:

Employee **AS** E(FN, M, LN, SSN, BD, Addr, Sex, Sal, SSSN, DNO)

Unspecified WHERE-clause

- ▶ A missing WHERE-clause indicates no condition: all tuples of the relations in the FROM-clause are selected
- ▶ This is equivalent to the condition WHERE TRUE

Query 4: Retrieve the SSN values for all employees

Q4: **SELECT** SSN
 FROM Employee;

Unspecified WHERE-clause

- ▶ If more than one relation is specified in the FROM-clause and there is no join condition, then the CARTESIAN PRODUCT of tuples is selected

Query 5: retrieve all combinations of Employee.SSN and Department.Dname

Q5: **SELECT** SSN, Dname
 FROM Employee, Department;

- ▶ It is extremely important not to overlook specifying any selection and join conditions in the WHERE-clause; otherwise, incorrect and very large relations may result

Use of ASTERISK (*)

- ▶ An asterisk (*) stands for *all the attributes*

Query 6: retrieves all the attribute values of any Employee who works in Department number 5

Q6: **SELECT** *
 FROM Employee
 WHERE DNO = 5;

Query 7: retrieves all the attributes of an Employee and the attributes of the Department in which he or she works for every employee of the 'Research' department

Q7: **SELECT** *
 FROM Employee, Department
 WHERE Dname = 'Research' **AND** DNO = Dnumber;

Use of DISTINCT

- ▶ SQL does not treat a relation as a set: duplicate tuples can appear in a query result.
- ▶ To eliminate duplicate tuples, use the keyword DISTINCT

Query 8: Retrieve the salary of every employee (Q8A) and all distinct salary values (Q8B)

Q8a: **SELECT** Salary
FROM Employee;

Q8b: **SELECT** **DISTINCT** Salary
FROM Employee;

- ▶ The result of Q8A may have duplicate SALARY values, but Q8B's

Examples

1. Retrieve the names of all employees in the departments which are located in Houston
2. List the names of all employees who have a dependent with the same first name as themselves
3. Make a list of all project numbers for projects that involve an employee whose last name is 'Smith' as a as a manager of the department that controls the project.

Set Operations

- ▶ Set union (**UNION**), set difference (**EXCEPT**) and set intersection (**INTERSECT**) operations
- ▶ The resulting relations of these set operations are sets of tuples: *duplicate tuples are eliminated from the result*
- ▶ The set operations apply only to *union compatible relations*
- ▶ UNION ALL, EXCEPT ALL, INTERSECT ALL

Set Operations

Query 9: Make a list of all project numbers for projects that involve an employee whose last name is 'Smith' as a worker or as a manager of the department that controls the project.

Q10: **(SELECT DISTINCT Pnumber**
FROM Project, Department, Employee
WHERE Dnum = Dnumber **AND** MgrSSN = SSN
AND Lname = 'Smith')
UNION
(SELECT DISTINCT Pnumber
FROM Works_on, Employee
WHERE ESSN=SSN **AND** Lname = 'Smith');

Substring pattern matching and arithmetic operators

- ▶ Two reserved characters: % and _

Query 10: Retrieve all employees whose address is in Houston, Texas.

Q10: SELECT *
FROM Employee
WHERE Address LIKE '%Houston,TX%';

Query 11: Retrieve all employees whose SSN has '88' at the end.

Q11: SELECT *
FROM Employee
WHERE SSN LIKE '_____88';

Substring pattern matching and arithmetic operators

- ▶ Standard arithmetic operators: +, -, *, /

Query 12: show the resulting salaries if every employee working on “ProductX” is given 10% raise

Q12: **SELECT** Fname, Lname, 1.1*Salary **AS** INC_SAL
FROM Employee, Works_on, Project
WHERE SSN = ESSN **AND** PNO = Pnumber
AND Pname = 'ProductX';

NULL & 3-valued logic

AND	True	False	Unknown
True	T	F	U
False	F	F	F
Unknown	U	F	U

OR	True	False	Unknown
True	T	T	T
False	T	F	U
Unknown	T	U	U

NOT	
True	F
False	T
Unknown	U

SELECT * FROM Employee WHERE SuperSSN IS NULL;

SELECT * FROM Employee WHERE SuperSSN IS NOT NULL;



SELECT Command

```
SELECT [DISTINCT | ALL]
      { * | [columnExpression [AS newName]] [, ...] }
FROM TableName [alias] [, ...]
[WHERE condition]
[GROUP BY columnList]    [HAVING condition]
[ORDER BY columnList]
```

Nested Queries

- ▶ Complete SELECT-FROM-WHERE blocks within WHERE clause of another query
- ▶ Comparison operator IN
 - ▶ Compares value v with a set (or multiset) of values V
 - ▶ Evaluates to *TRUE* if v is one of the elements in V

Query 13: Retrieve the name and address of all employees who work for the 'Research' department

Q13: **SELECT** Fname, Lname, Address
 FROM Employee
 WHERE Dno **IN** (**SELECT** Dnumber
 FROM Department
 WHERE Dname = 'Research');

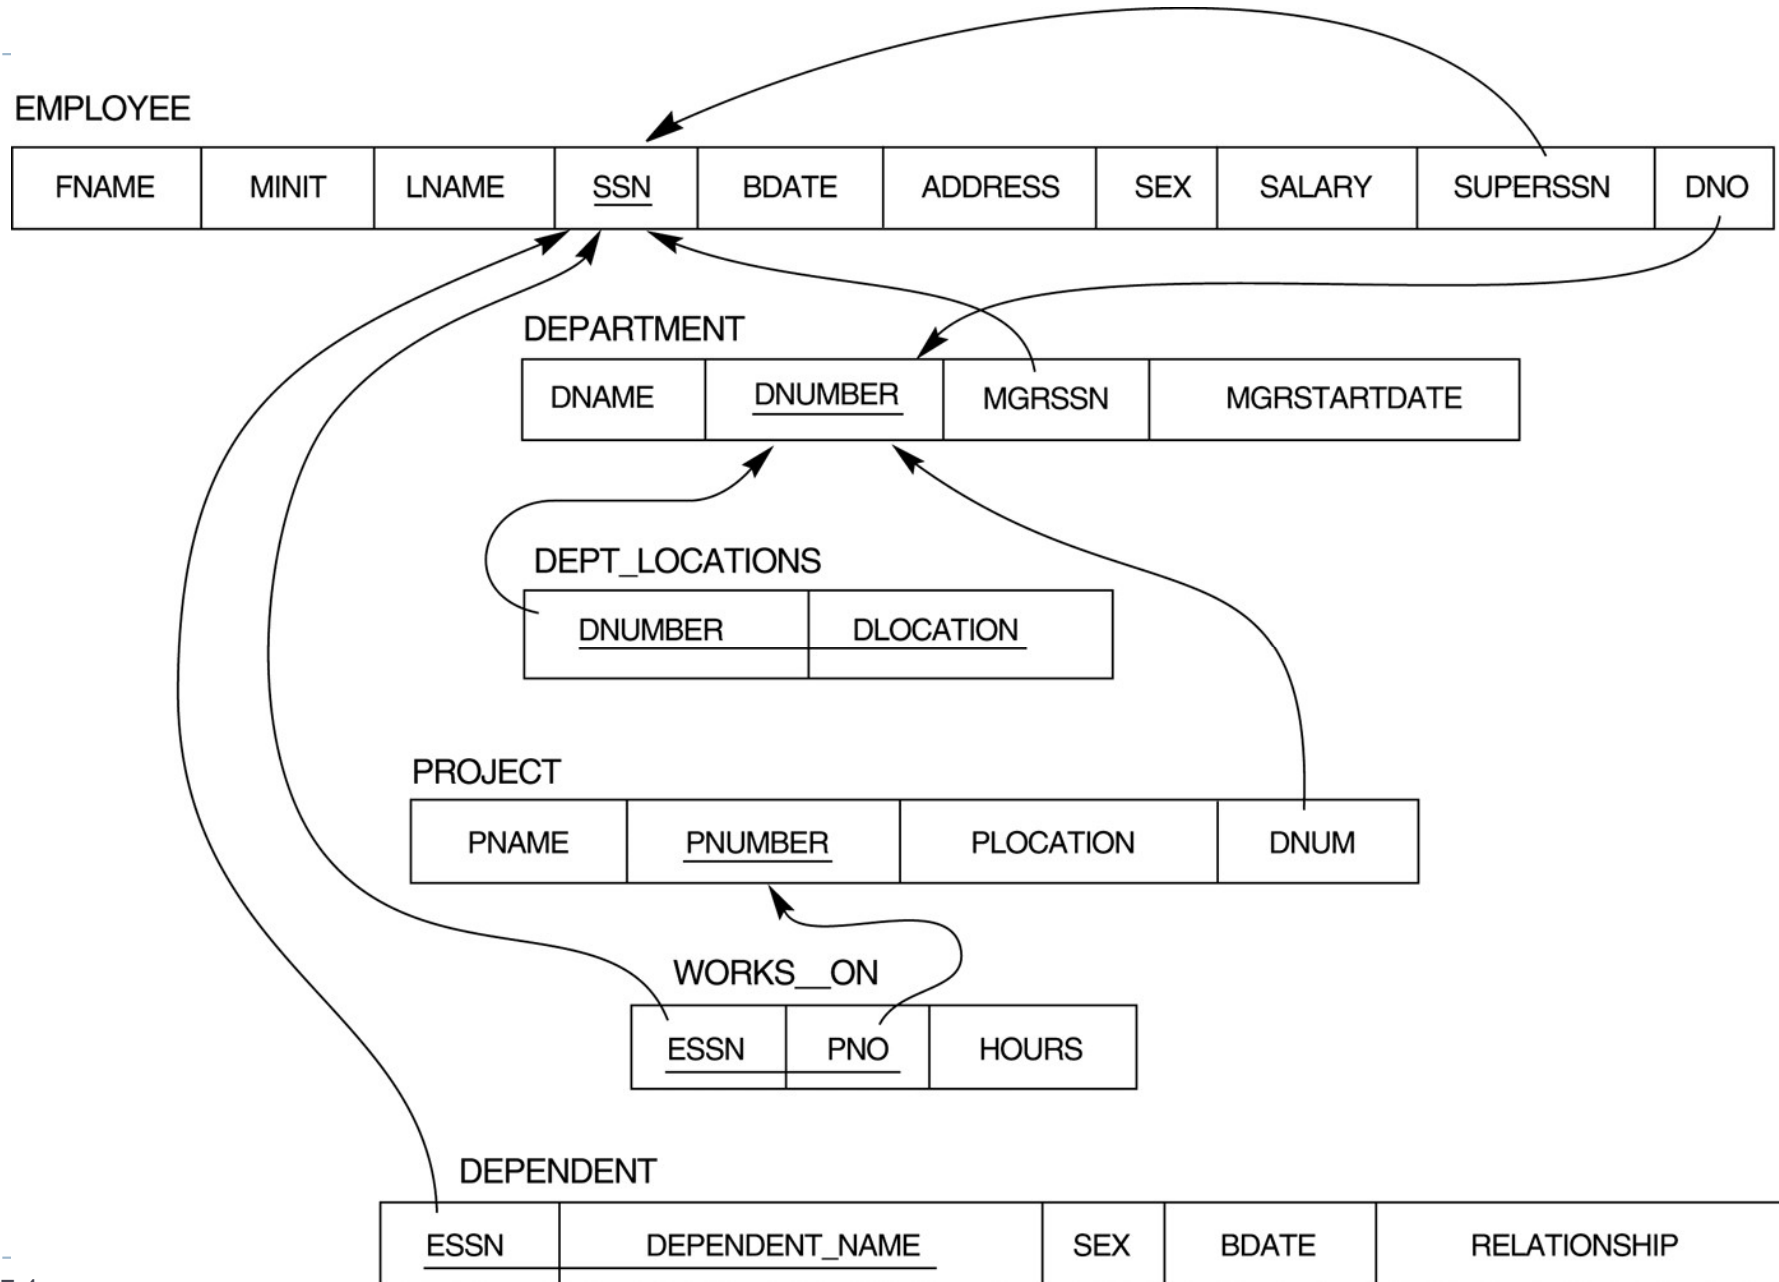
Correlated Nested Queries

- ▶ If a condition in the WHERE-clause of a nested query references an attribute of a relation declared in the outer query, the two queries are said to be **correlated**

Query 14: Retrieve the name of each employee who has a dependent with the same first name as the employee.

Q14: **SELECT** E.Fname, E.Lname
FROM Employee **E**
WHERE E.SSN **IN** (**SELECT** ESSN
FROM Dependent
WHERE ESSN = **E.SSN** **AND**
E.Fname = Dependent_name);

The COMPANY Database



Correlated Nested Queries

- ▶ A query written with nested SELECT-FROM-WHERE blocks and using IN comparison operator can always be expressed as a single block query
- ▶ For example, Q14 may be written as in Q14A

Q14a: **SELECT** E.Fname, E.Lname
FROM Employee E, Dependent D
WHERE E.SSN = D.ESSN **AND**
E.Fname = D.Dependent_name;

Nested Query Exercises

Query 15: Retrieve the SSNs of all employees who work the same (project, hours) combination on some project that employee John Smith (SSN=123456789) works on (using a nested query)

**Q15: SELECT DISTINCT ESSN
FROM Works_on
WHERE (PNO, Hours) IN
(SELECT PNO, Hours
FROM Works_on
WHERE ESSN = '123456789');**

More Comparison Operators

- ▶ Operators that can be combined with ANY (or SOME), ALL: =, >, >=, <, <=, and <>

Query 16: Retrieve all employees whose salary is greater than the salary of all employees in department 5

**Q16: SELECT *
FROM Employee
WHERE Salary > ALL (SELECT Salary
FROM Employee
WHERE DNO=5);**

EXISTS and UNIQUE Functions

- ▶ **EXISTS and NOT EXISTS** function
 - ▶ Typically used in conjunction with a correlated nested query
 - ▶ EXISTS(Q) returns TRUE if the result of a query Q is NOT empty (Some tuples EXIST in the result).
 - ▶ NOT EXISTS(Q) returns TRUE if the result of a query Q is empty (No tuples are in the result).
- ▶ **UNIQUE(Q)** function
 - ▶ Returns TRUE if there are no duplicate tuples in the result of query Q

EXISTS Function

Query 14: Retrieve the name of each employee who has a dependent with the same first name as the employee

Q14b: **SELECT** Fname, Lname
FROM Employee
WHERE EXISTS (SELECT *
FROM Dependent
WHERE ESSN = SSN AND
FName = Dependent_name);

EXISTS Function

Query 17: Retrieve the names of employees who have no dependents

**Q17: SELECT Fname, Lname
FROM Employee
WHERE NOT EXISTS (SELECT *
FROM Dependent
WHERE SSN = ESSN);**

- ▶ In Q17, the correlated nested query retrieves all DEPENDENT tuples related to an EMPLOYEE tuple. If none exist, the EMPLOYEE tuple is selected

Enumerated Sets

- ▶ An explicit (enumerated) set of values in the WHERE-clause

Query 18: Retrieve the SSNs of all employees who work on project numbers 1, 2, or 3.

**Q18: SELECT DISTINCT ESSN
FROM Works_on
WHERE PNO IN (1, 2, 3);**

Joined Relations

- ▶ Can specify a "joined relation" in the FROM-clause
- ▶ Allows the user to specify different types of joins
 - ▶ EQUIJOIN
 - ▶ NATURAL JOIN
 - ▶ LEFT OUTER JOIN
 - ▶ RIGHT OUTER JOIN
 - ▶ FULL OUTER JOIN

Joined Tables and Outer Joins

- ▶ Joined table
 - ▶ Permits users to specify a table resulting from a join operation in the FROM clause of a query

Query 1: Retrieve the name and address of all employees who work for the 'Research' department.

Q1a: **SELECT** Fname, Lname, Address
 FROM (Employee **JOIN** Department **ON** Dno = Dnumber)
 WHERE Dname = 'Research';

Q1: **SELECT** Fname, Lname, Address
 FROM Employee, Department
 WHERE Dname='Research' **AND** Dnumber= Dno;

Joined Tables and Outer Joins

- ▶ Specify different types of join
 - ▶ NATURAL JOIN
 - ▶ Various types of OUTER JOIN
- ▶ NATURAL JOIN on two relations R and S
 - ▶ **No join condition specified**
 - ▶ Implicit EQUIJOIN condition for each pair of attributes with same name from R and S

Joined Tables in SQL and Outer Joins (cont'd.)

▶ **Inner join**

- ▶ Default type of join in a joined table
- ▶ Tuple is included in the result only if a matching tuple exists in the other relation

▶ **LEFT OUTER JOIN**

- ▶ Every tuple in LEFT table must appear in result
- ▶ If no matching tuple
 - Padded with NULL values for attributes of RIGHTtable

Joined Tables in SQL and Outer Joins (cont'd.)

- ▶ RIGHT OUTER JOIN
 - ▶ Every tuple in RIGHT table must appear in result
 - ▶ If no matching tuple
 - Padded with NULL values for the attributes of LEFT table
- ▶ FULL OUTER JOIN

Joined Relations - Examples

Query 3: For each employee, retrieve the employee's name, and the name of his or her immediate supervisor.

Q3a: **SELECT** E.Fname, E.Lname, S.Fname, S.Lname
FROM Employee E S
WHERE E.SuperSSN = S.SSN;

Q3c: **SELECT** E.Fname, E.Lname, S.Fname, S.Lname
FROM (Employee E **LEFT OUTER JOIN**
Employee S **ON** E.SuperSSN = S.SSN);

► **Compare two queries???**

Joined Relations - Examples

Query 1: Retrieve the name and address of all employees who work for the 'Research' department.

Q1: **SELECT** Fname, Lname, Address
 FROM Employee, Department
 WHERE Dname = 'Research' **AND** Dnumber = Dno;

► could be written as:

Q1a: **SELECT** Fname, Lname, Address
 FROM (Employee **JOIN** Department **ON** Dnumber = Dno)
 WHERE Dname = 'Research';

Q1b: **SELECT** Fname, Lname, Address
 FROM (Employee **NATURAL JOIN** (Department
 AS Dept(Dname, Dno, MSSN, MSDate)))
 WHERE Dname = 'Research';

Joined Relations - Examples

Query 2: For every project located in 'Stafford', list the project number, the controlling department number, and the department manager's last name, address, and birthdate

Q2a: **SELECT** Pnumber, Dnum, Lname, Bdate, Address
FROM ((Project **JOIN** Department **ON** Dnum =
Dnumber) **JOIN** Employee **ON** MGRSSN = SSN))
WHERE Plocation = 'Stafford' ;

AGGREGATE FUNCTIONS

► COUNT, SUM, MAX, MIN, AVG

Query 19: Find the max, min, & average salary among all employees

**Q19: SELECT MAX(Salary), MIN(Salary), AVG(Salary)
FROM Employee;**

AGGREGATE FUNCTIONS

Queries 20: Retrieve the total number of employees in the company

**Q20: SELECT COUNT (*)
FROM Employee;**

Queries 21: Retrieve the number of employees in the 'Research' department

**Q21: SELECT COUNT (*)
FROM Employee, Department
WHERE Dno = Dnumber AND Dname = 'Research';**

- ▶ Note: NULL values are discarded wrt. aggregate functions as applied to a particular column

GROUPING

- ▶ A GROUP BY-clause is for specifying the grouping attributes, which must also appear in the SELECT-clause
- ▶ Each subgroup of tuples consists of the set of tuples that have the same value for the grouping attribute(s)
- ▶ Apply the aggregate functions to subgroups of tuples in a relation
- ▶ Each subgroup of tuples consists of the set of tuples that have the same value for the grouping attribute(s)
- ▶ The aggregate function is applied to each subgroup independently
- ▶ If NULLs exist in grouping attribute
 - ▶ Separate group created for all tuples with a NULL value in grouping attribute

SELECT Command

```
SELECT [DISTINCT | ALL]
      { * | [columnExpression [AS newName]] [,...] }
FROM TableName [alias] [, ...]
[WHERE condition]
[GROUP BY columnList]    [HAVING condition]
[ORDER BY columnList]
```

GROUPING

Query 22: For each department, retrieve the department number, the number of employees in the department, and their average salary

Q22: **SELECT** Dno, **COUNT** (*), **AVG** (Salary)
 FROM Employee
 GROUP BY Dno;

- ▶ In Q22, the EMPLOYEE tuples are divided into groups - each group having the same value for the grouping attribute DNO
- ▶ The COUNT and AVG functions are applied to each such group of tuples separately
- ▶ The SELECT-clause includes only the grouping attribute and the functions to be applied on each group of tuples
- ▶ A join condition can be used in conjunction with grouping

GROUPING: Q22 result

FNAME	MINIT	LNAME	<u>SSN</u>	• • •	SALARY	SUPERSSN	DNO
John	B	Smith	123456789	• • •	30000	333445555	5
Franklin		Wong	333445555		40000	888665555	5
Ramesh	K	Narayan	666884444		38000	333445555	5
Joyce	A	English	453453453		25000	333445555	5
Alicia	J	Zelaya	999887777		25000	987654321	4
Jennifer	S	Wallace	987654321		43000	888665555	4
Ahmad	V	Jabbar	987987987		25000	987654321	4
James	E	Bong	888665555		55000	null	1

DNO	COUNT (*)	AVG (SALARY)
5	4	33250
4	3	31000
1	1	55000

Result of Q22

Grouping EMPLOYEE tuples by the value of DNO.

GROUPING: THE HAVING-CLAUSE

- ▶ Sometimes we want to retrieve the values of these functions for only those groups that satisfy certain conditions
- ▶ **The HAVING-clause** is used for specifying a **selection condition on groups** (rather than on individual tuples)

GROUPING: THE HAVING-CLAUSE

Query 23: For each project on which more than two employees work , retrieve the project number, project name, and the number of employees who work on that project.

Q23: **SELECT** Pnumber, Pname, **COUNT** (*)
 FROM Project, Works_on
 WHERE Pnumber = Pno
 GROUP BY Pnumber, Pname
 HAVING **COUNT** (*) > 2;

ORDER BY

- ▶ The ORDER BY clause is used to sort the tuples in a query result based on the values of some attribute(s)

Query 24: Retrieve a list of employees and the projects each works in, ordered by the employee's department, and within each department ordered alphabetically by employee last name

Q24: **SELECT** Dname, Lname, Fname, Pname
FROM Department, Employee, Works_on, Project
WHERE Dnumber = Dno **AND** SSN = ESSN
AND Pno = Pnumber
ORDER BY Dname, Lname [DESC|ASC]

SELECT Command

```
SELECT [DISTINCT | ALL]
      { * | [columnExpression [AS newName]] [, ...] }
FROM TableName [alias] [, ...]
[WHERE condition]
[GROUP BY columnList]    [HAVING condition]
[ORDER BY columnList]
```

SELECT Command

- ▶ **SELECT** Specifies which columns are to appear in output
- ▶ **FROM** Specifies table(s) to be used
- ▶ **WHERE** Filters rows
- ▶ **GROUP BY** Forms groups of rows with same column value
- ▶ **HAVING** Filters groups subject to some condition
- ▶ **ORDER BY** Specifies the order of the output

Contents

-
- 1 Introduction of Structured Query Language
 - 2 DDL: create, drop, alter
 - 3 DML: select, insert, update, delete**
 - 4 DCL: commit, rollback, grant, revoke

Insert Command

- ▶ Add one or more tuples to a relation
- ▶ Attribute values should be listed in the same order as the attributes were specified in the CREATE TABLE command

```
INSERT INTO TableName (Attribute1, Attribute2, ...)
VALUES (value1, value2, ...);
```

Insert Command

- ▶ Insert a tuple for a new EMPLOYEE:

U1: **INSERT INTO** Employee
 VALUES ('Richard', 'K', 'Marini', '653298653',
 '30-DEC-52', '98 Oak Forest, Katy, TX',
 'M', 37000, '987654321', 4);

- ▶ An alternate form of INSERT specifies explicitly the attribute names that correspond to the values in the new tuple, attributes with NULL values can be left out
- ▶ Example: Insert a tuple for a new EMPLOYEE for whom we only know the FNAME, LNAME, and SSN attributes.

U2: **INSERT INTO** Employee (Fname, Lname, SSN)
 VALUES ('Richard', 'Marini', '653298653');

Insert Command

- ▶ Important note: Only the constraints specified in the DDL commands are automatically enforced by the DBMS when updates are applied to the database
- ▶ Another variation of INSERT allows insertion of multiple tuples resulting from a query into a relation

Insert Command

- ▶ Example: *Suppose we want to create a temporary table that has the name, number of employees, and total salaries for each department. A table DEPTS_INFO is created by U3, and is loaded with the summary information retrieved from the database by the query in U3A*

U3: **CREATE TABLE** Depts_info
 (Dept_name VARCHAR(10),
 No_of_emps INTEGER,
 Total_sal INTEGER);

U3A: **INSERT INTO** Depts_info (Dept_name, No_of_emps,
 Total_sal)
 SELECT Dname, **COUNT** (*), **SUM** (Salary)
 FROM Department, Employee
 WHERE Dnumber = Dno
 GROUP BY Dname;

Delete Command

DELETE FROM *TableName*
WHERE *Condition*;

- ▶ Removes tuples from a relation
- ▶ Tuples are deleted from only one table at a time (unless CASCADE is specified on a referential integrity constraint)
- ▶ A missing WHERE-clause specifies that all tuples in the relation are to be deleted; the table then becomes an empty table
- ▶ The number of tuples deleted depends on the number of tuples in the relation that satisfy the WHERE-clause

Delete Command - Examples

U4A: DELETE FROM Employee
WHERE Lname = 'Brown';

U4B: DELETE FROM Employee
WHERE SSN = '123456789';

U4C: DELETE FROM Employee
WHERE Dno IN
(SELECT Dnumber
FROM Department
WHERE Dname = 'Research');

U4D: DELETE FROM Employee;

Update Command

UPDATE *TableName*

SET *Set-Clause*

WHERE *Condition;*

- ▶ Used to modify attribute values of one or more selected tuples
- ▶ A WHERE-clause selects the tuples to be modified
- ▶ An additional SET-clause specifies the attributes to be modified and their new values
- ▶ Each command modifies tuples in the same relation
- ▶ Referential integrity should be enforced

Update Command

- ▶ Example: *Change the location and controlling department number of project number 10 to 'Bellaire' and 5, respectively.*

U5: **UPDATE** Project
 SET Plocation = 'Bellaire', Dnum = 5
 WHERE Pnumber = 10;

Update Command

- ▶ Example: *Give all employees in the 'Research' department a 10% raise in salary.*

U6: **UPDATE** Employee
 SET Salary = Salary *1.1
 WHERE Dno **IN** (**SELECT** Dnumber
 FROM Department
 WHERE Dname = 'Research');

Advanced DDL: Assertions & Triggers

- ▶ ASSERTIONs to express constraints that do not fit in the basic SQL categories
- ▶ Mechanism: CREATE ASSERTION
 - ▶ components include: a constraint name, followed by CHECK, followed by a condition

Advanced DDL: Assertions & Triggers

- ▶ Example: *The salary of an employee must not be greater than the salary of the manager of the department that the employee works for*

CREATE ASSERTION Salary_constraint

CHECK (NOT EXISTS (SELECT *

FROM Employee E, Employee M,
Department D

WHERE E.Salary > M.Salary **AND**
E.Dno = D.Number **AND**
D.MGRSSN = M.SSN));

Advanced DDL: Assertions & Triggers

- ▶ Triggers: to specify the type of action to be taken as certain events occur and as certain conditions are satisfied
- ▶ Details of triggers: presentation and lab

Views

- ▶ A view is a “virtual” table that is derived from other tables
- ▶ Allows for limited update operations (since the table may not physically be stored)
- ▶ Allows full query operations
- ▶ A convenience for expressing certain operations

VIEWS

- Specify a different WORKS_ON table (view)

```
CREATE VIEW      Works_on_new AS  
      SELECT      Fname, Lname, Pname, Hours  
      FROM        Employee, Project, Works_on  
      WHERE       SSN = ESSN AND Pno = Pnumber;
```

- We can specify SQL queries on a newly create table (view):

```
SELECT Fname, Lname From Works_on_new  
WHERE Pname = 'Seena';
```

- When no longer needed, a view can be dropped:

```
DROP VIEW Works_on_new;
```

View Update and Inline Views

- ▶ Update on a view defined on a single table without any aggregate functions
 - ▶ Can be mapped to an update on underlying base table
- ▶ View involving joins
 - ▶ Often not possible for DBMS to determine which of the updates is intended

Contents

-
- 1 Introduction of Structured Query Language
 - 2 DDL: create, drop, alter
 - 3 DML: select, insert, update, delete
 - 4 DCL: commit, rollback, grant, revoke
(Chapter 9 + homework)**

SQL for Data Control

- ▶ Commands:
 - ▶ GRANT
 - ▶ REVOKE
- ▶ Based on three central objects:
 - ▶ Users
 - ▶ Database objects
 - ▶ Privileges: select, modify (insert, update, delete), reference



SQL for Data Control

- ▶ GRANT: pass privileges on their own database objects to other users

GRANT <privilege list>
ON <database objects>
TO <user list>

- ▶ REVOKE: take back (cancel) privileges on their own database objects from other users

REVOKE <privilege list>
ON <database objects>
FROM <user list>



SQL for Data Control

- ▶ Propagation of Privileges using the GRANT OPTION
 - ▶ Whenever the owner A of a relation R grants a privilege on R to another account B, privilege can be given to B with or without the GRANT OPTION.
 - ▶ If the GRANT OPTION is given, this means that B can also grant that privilege on R to other accounts.



An Example

- ▶ Suppose that the DBA creates four accounts
 - ▶ A1, A2, A3, A4
- ▶ and wants only A1 to be able to create base relations. Then the DBA must issue the following GRANT command in SQL

GRANT CREATETAB TO A1;

- ▶ In SQL2 the same effect can be accomplished by having the DBA issue a **CREATE SCHEMA** command as follows:

CREATE SCHEMA EXAMPLE **AUTHORIZATION**
A1;



An Example(2)

- ▶ User account A1 can create tables under the schema called **EXAMPLE**.
- ▶ Suppose that A1 **creates** the two base relations **EMPLOYEE** and **DEPARTMENT**
 - ▶ A1 is then **owner** of these two relations and hence all the relation privileges on each of them.
- ▶ Suppose that A1 wants to grant A2 the privilege to insert and delete tuples in both of these relations, but A1 does not want A2 to be able to propagate these privileges to additional accounts:
GRANT INSERT, DELETE ON
EMPLOYEE, DEPARTMENT TO A2;



An Example(3)

EMPLOYEE

Name	<u>Ssn</u>	Bdate	Address	Sex	Salary	Dno
------	------------	-------	---------	-----	--------	-----

DEPARTMENT

<u>Dnumber</u>	Dname	Mgr_ssn
----------------	-------	---------

Figure 23.1

Schemas for the two relations EMPLOYEE and DEPARTMENT.



An Example(4)

- ▶ Suppose that A1 wants to allow A3 to retrieve information from either of the two tables and also to be able to propagate the SELECT privilege to other accounts.
 - ▶ A1 can issue the command:
**GRANT SELECT ON EMPLOYEE, DEPARTMENT
TO A3 WITH GRANT OPTION;**
 - ▶ A3 can grant the **SELECT** privilege on the **EMPLOYEE** relation to A4 by issuing:
GRANT SELECT ON EMPLOYEE TO A4;
 - ▶ Notice that A4 can't propagate the SELECT privilege because GRANT OPTION was not given to A4
-



An Example(5)

- ▶ Suppose that A1 decides to revoke the SELECT privilege on the EMPLOYEE relation from A3; A1 can issue:

REVOKE SELECT ON EMPLOYEE FROM A3 ;

- ▶ The DBMS must now automatically revoke the SELECT privilege on EMPLOYEE from A4, too, because A3 granted that privilege to A4 and A3 does not have the privilege any more.



An Example(6)

- ▶ Suppose that A1 wants to give back to A3 a limited capability to **SELECT** from the **EMPLOYEE** relation and wants to allow A3 to be able to propagate the privilege.
 - ▶ The limitation is to retrieve only the **NAME**, **BDATE**, and **ADDRESS** attributes and only for the tuples with **DNO=5**.

- ▶ A1 then create the view:

```
CREATE VIEW A3EMPLOYEE AS  
  SELECT NAME, BDATE, ADDRESS  
  FROM EMPLOYEE  
  WHERE DNO = 5;
```

- ▶ After the view is created, A1 can grant **SELECT** on the view **A3EMPLOYEE** to A3 as follows:

```
GRANT SELECT ON A3EMPLOYEE TO A3  
  WITH GRANT OPTION;
```



An Example(7)

- ▶ Finally, suppose that A1 wants to allow A4 to update only the SALARY attribute of EMPLOYEE;
- ▶ A1 can issue:
GRANT UPDATE ON EMPLOYEE (SALARY) TO A4 ;
- ▶ The UPDATE or INSERT privilege can specify particular attributes that may be updated or inserted in a relation.
 - ▶ Other privileges (SELECT, DELETE) are not attribute specific.



Summary

- ▶ SQL developments: an overview
- ▶ SQL
 - ▶ DDL: Create, Alter, Drop
 - ▶ DML: select, insert, update, delete
 - ▶ Introduction to advanced DDL (assertions & triggers), views, DCL (commit, rollback, grant, revoke)



Exercise

EMPLOYEE

FNAME	MINIT	LNAME	<u>SSN</u>	BDATE	ADDRESS	SEX	SALARY	SUPERSSN	DNO
-------	-------	-------	------------	-------	---------	-----	--------	----------	-----

DEPARTMENT

DNAME	<u>DNUMBER</u>	MGRSSN	MGRSTARTDATE	CREATEDATE
-------	----------------	--------	--------------	------------

EMPLOYEE:

- ▶ Fname, Lname: VARCHAR(15), NOT NULL
- ▶ Minit: CHAR
- ▶ SSN: CHAR(9), NOT NULL, PRIMARY KEY
- ▶ Bdate: DATE, <= "1/1/1999"
- ▶ Address: VARCHAR(100)
- ▶ Sex: CHAR, {F/M}
- ▶ Salary: DECIMAL(10,2)
- ▶ SuperSSN: CHAR(9), refers to EMPLOYEE(SSN)
- ▶ Dno: INT, NOT NULL, default value = 1, refers to DEPARTMENT(Dnumber) –
ON DELETE SET DEFAULT

DEPARTMENT:

- ▶ Dname: VARCHAR(15), NOT NULL, UNIQUE
- ▶ Dnumber: INT, NOT NULL, PRIMARY KEY
- ▶ MgrSSN: CHAR(9), NOT NULL, default value = '888665555', refers to EMPLOYEE(SSN) – ON DELETE SET DEFAULT, ON UPDATE CASCADE
- ▶ MgrStartDate: DATE
- ▶ CreateDate: DATE, <= MgrStartDate

EMPLOYEE

Fname	Minit	Lname	SSN	Bdate	Address	Sex	Salary	SuperSSN	DNO
An	V	Nguyen	123456789	1/1/1990	TP.HCM	F	10,000	987654321	1
Binh	T	Nguyen	987654321	2/2/1988	Ha noi	M	15,000		1
Hoa	T	Tran	111222333	3/3/1991	Binh Duong	F	12,000	123456789	2
Long	K	Ly	888665555	4/4/1993	Dong Nai	M	20,000	987654321	2

DEPARTMENT

Dname	Dnumber	MgrSSN	MgrStartDate	CreateDate
NH	1	123456789	10/10/2010	1/1/1999
TC	2	888665555	5/5/2000	1/1/1999

CREATE TABLE

```
CREATE TABLE [SchemaName.]TableName
({colName dataType [NOT NULL] [UNIQUE] [PRIMARY KEY]
[DEFAULT defaultOption]
[CHECK searchCondition] [...]}
[PRIMARY KEY (listOfColumns),]
{[UNIQUE (listOfColumns),] [...]}
{[FOREIGN KEY (listOfFKColumns)
REFERENCES ParentTableName [(listOfCKColumns)]
[ON UPDATE referentialAction]
[ON DELETE referentialAction] [...]}
{[CHECK (searchCondition)] [...] })
```

```
CREATE DEPARTMENT (  
Dname VARCHAR(15) NOT NULL UNIQUE,  
Dnumber INT NOT NULL PRIMARY KEY,  
MgrSSN: CHAR(9) NOT NULL DEFAULT  
'888665555',  
MgrStartDate DATE,  
CreateDate DATE,  
CHECK (CreatDate <= MgrStartDate)
```

EMPLOYEE

Fname	Minit	Lname	<u>Ssn</u>	Bdate	Address	Sex	Salary	Super_ssn	Dno
-------	-------	-------	------------	-------	---------	-----	--------	-----------	-----

DEPARTMENT

Dname	<u>Dnumber</u>	Mgr_ssn	Mgr_start_date
-------	----------------	---------	----------------

DEPT_LOCATIONS

Dnumber	Dlocation
---------	-----------

PROJECT

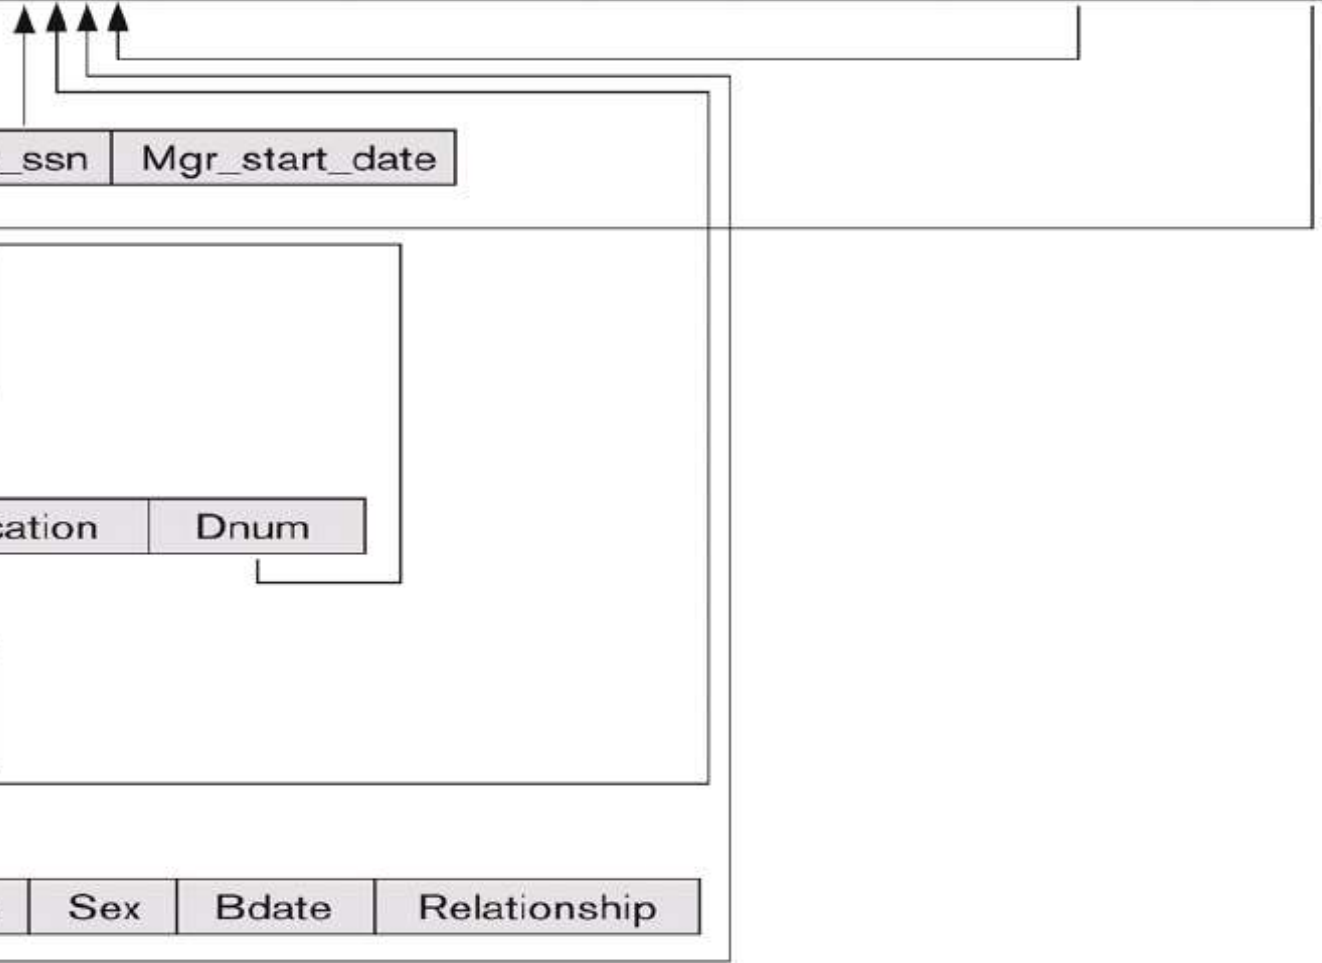
Pname	<u>Pnumber</u>	Plocation	Dnum
-------	----------------	-----------	------

WORKS_ON

<u>Essn</u>	<u>Pno</u>	Hours
-------------	------------	-------

DEPENDENT

<u>Essn</u>	<u>Dependent_name</u>	Sex	Bdate	Relationship
-------------	-----------------------	-----	-------	--------------



Exercise 2

1. Retrieve the names of all employees in the departments which are located in Houston
2. List the names of all employees who have a dependent with the same first name as themselves
3. For each project, calculate the total number of employees who work for it, and the total number of hours that these employees work for the project.
4. Retrieve the average salary of all female employees.
5. For each department whose average employee salary is more than \$30.000, retrieve the department name and the number of employees work for that department.

Exercise 3

1. Retrieve the name and address of all employees who work for the department which is managed by John B. Smith.
2. For every project located in 'Stafford', list the project number, the controlling department name, and the number of employees working on it (project).
3. Find the names of employees who work on all the projects controlled by department number 5.
4. List the names of all employees with more than 2 dependents.
5. Retrieve the names of employees who have no sons as dependents.
6. Retrieve the average salary of all female employees.
7. For each department whose average employee salary is more than \$30.000, retrieve the department name and the number of employees work for that department.